


“Spooky action at a distance”

Popular programming languages have converged on a common mutation model, where types behave either as **primitive/built-in** types or as **aggregate** types.


```
class Rect { Vec2 pos, dim; }
class Vec2 { Int x, y; }
```

```
var v1 = Vec2(2, 5);
var r1 = Rect(v1, v1);
var r2 = r1;
```



A variable of primitive type always has an **independent** value, but a variable of aggregate type has a value that is **shared** when assigned to another variable.

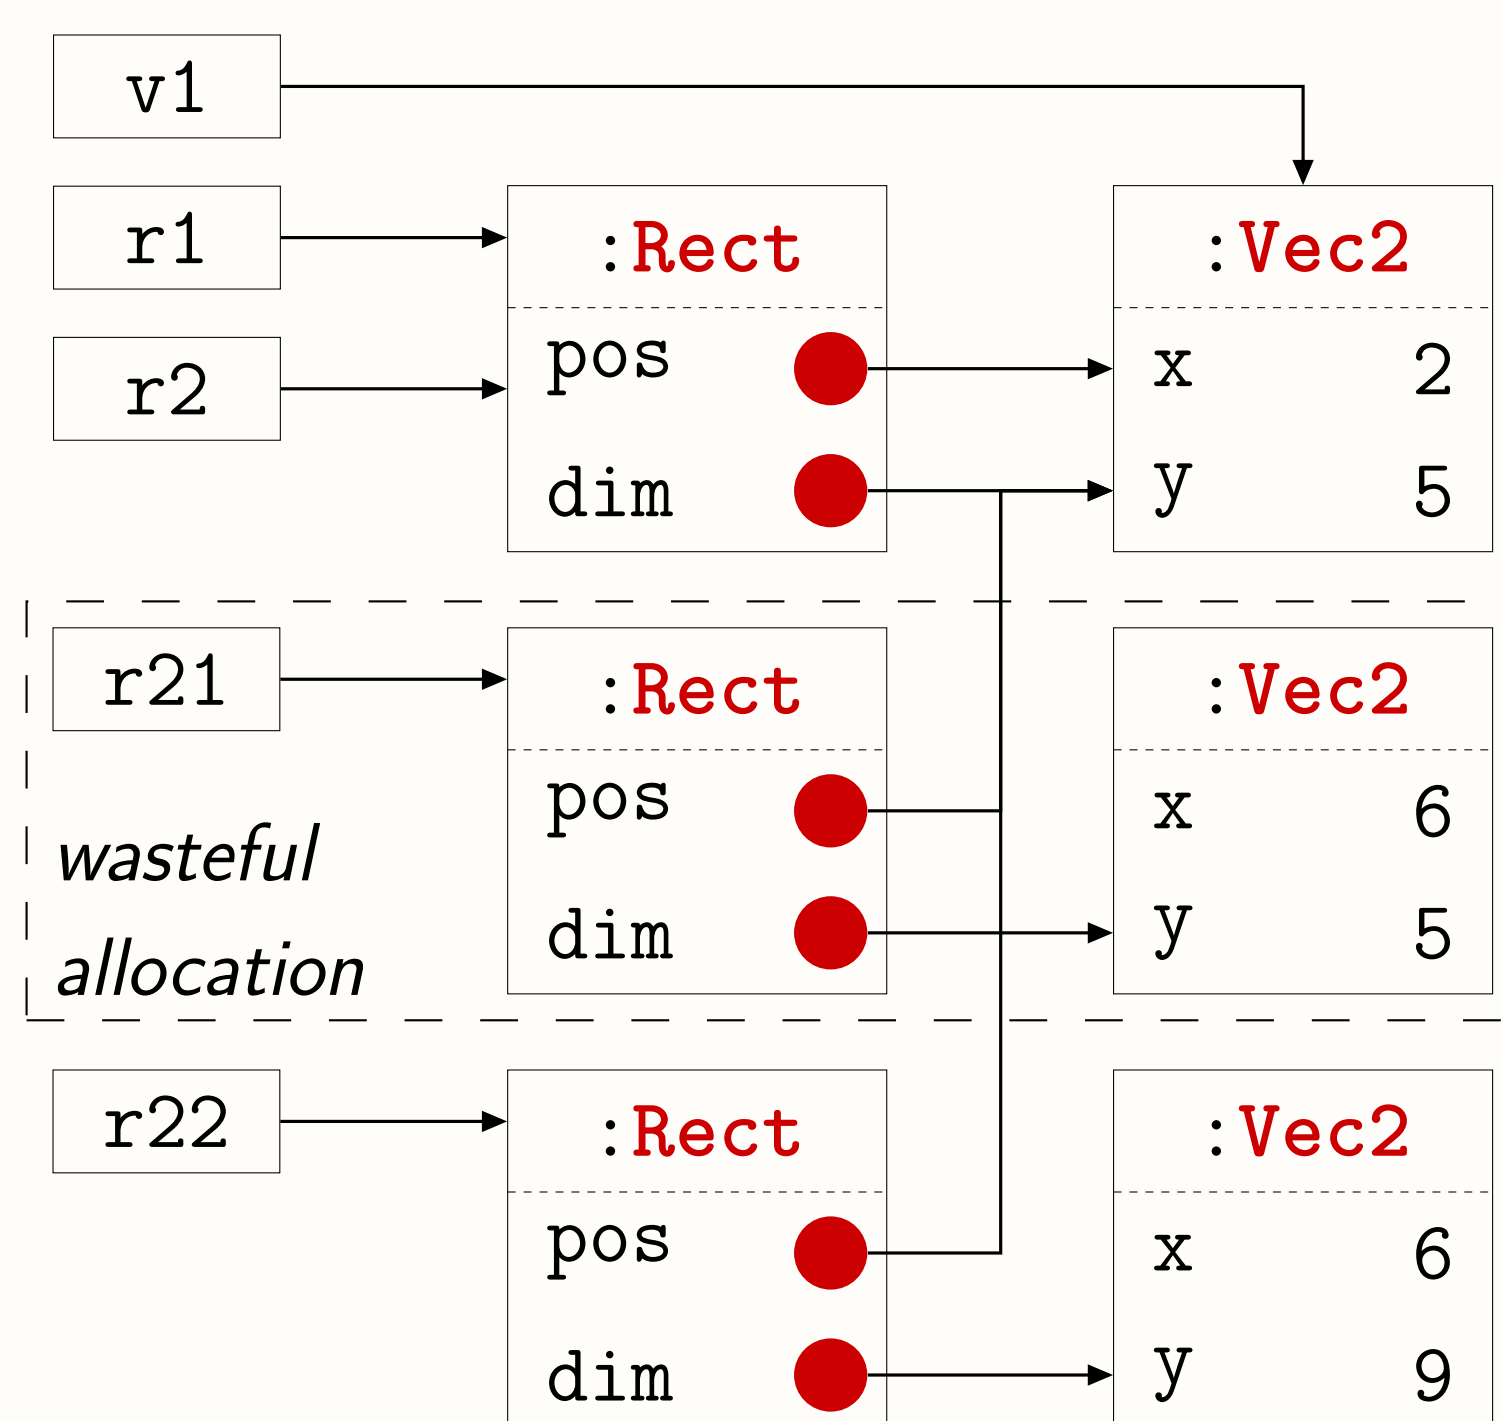
```
r2.dim.x += 4
r2.dim.y += 4
print(r1.pos.x) // Prints 6
```



Pure functional programming

Immutability may fail to capture the programmer's mental model, or prove ill-suited to express and optimize some algorithms. By design, functional updates do not support in-place mutation, demanding optimizer heroics to recover efficiency.

```
var r21 = up_dim(r2,
  up_x(r2.dim, r2.dim.x + 4))
var r22 = up_dim(r21,
  up_y(r21.dim, r21.dim.y + 4))
```

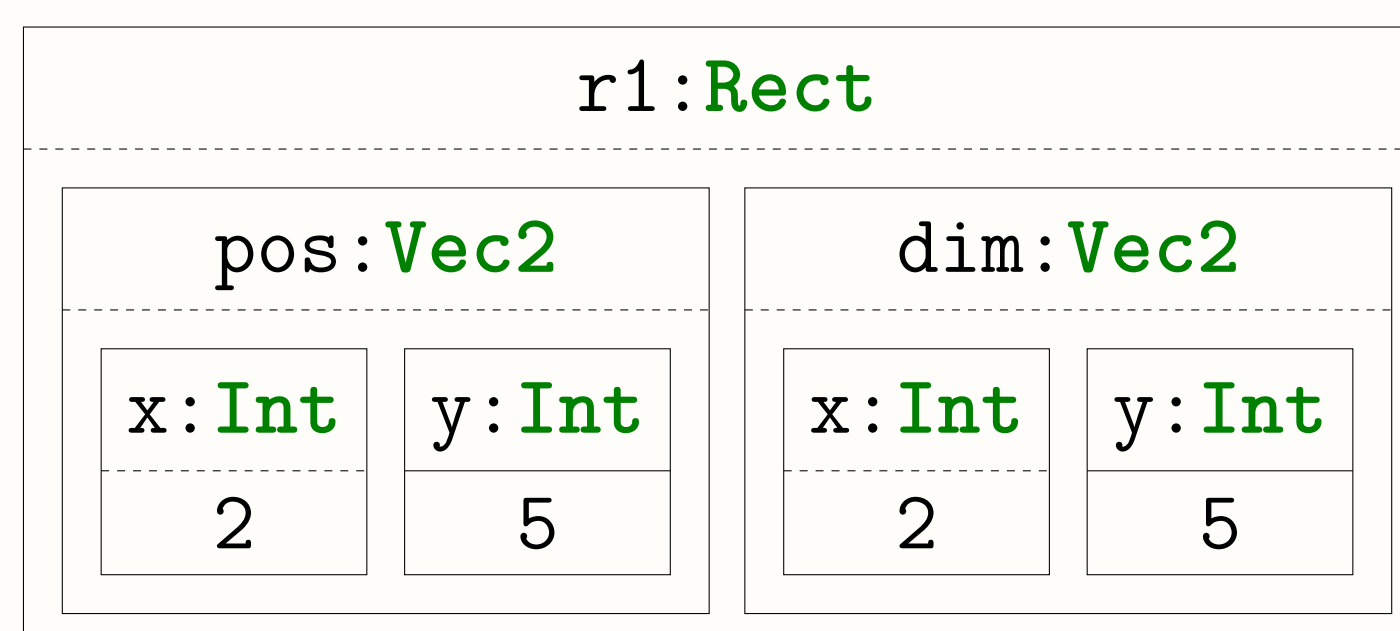


Advanced type systems

Newer programming languages (e.g., Rust) have blended different ideas aimed at taming aliasing into flow-sensitive type systems for more expressiveness. Unfortunately, these approaches have complexity costs that significantly raise the entry barrier for inexperienced developers [3].

Mutable value semantics

Mutable value semantics bans sharing instead of mutation to support in-place mutation and local reasoning. Conceptually, values are nested boxes, rather than graphs, which are copied when assigned or passed as parameters.



By design, all values are independent: mutating the part of a value has no impact on the value of other variables. Immutability and copy apply transitively, obviating the “shallow” vs “deep” distinction.

Mutation across function boundaries

Values of parameters whose type is annotated by **inout** are written back to their original location at the end of a function call.

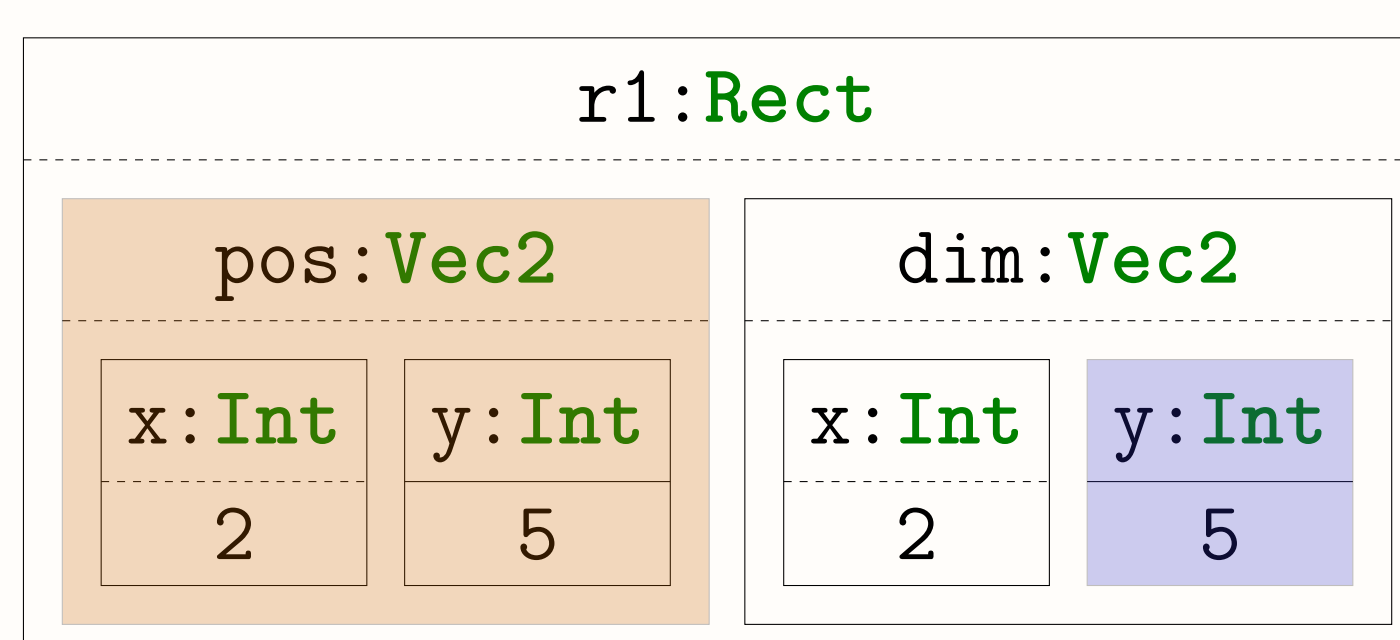
```
void scale_up(inout Vec2 v, Int f) {
  v.x *= f
  v.y *= f
}
```

```
scale_up(&r1.dim, 8)
print(r1.dim.x) // Prints 16
```

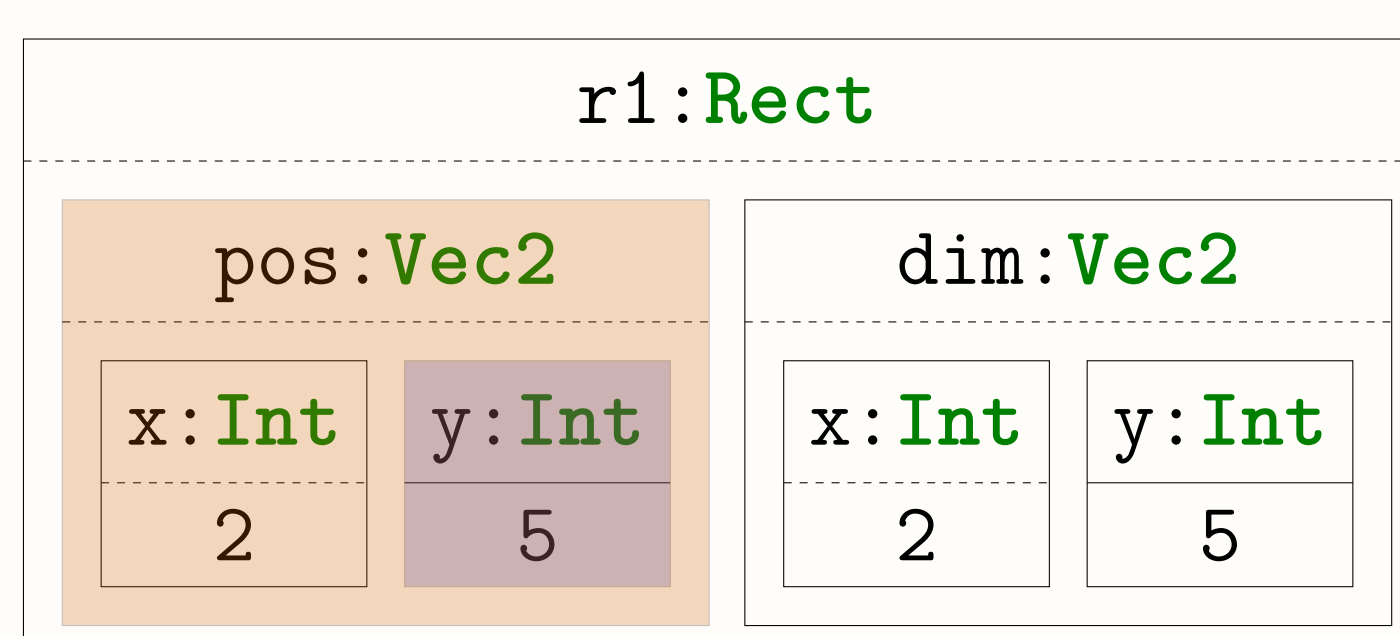
Overlapping mutations are prohibited to prevent any write-back from being discarded.

```
void swap_gt(inout Vec2 v,
  inout Int x) {...}
```

```
// Well-typed situation.
swap_gt(&r1.pos, &r1.dim.y)
```



```
// Ill-typed situation.
swap_gt(&r1.pos, &r1.pos.y)
```



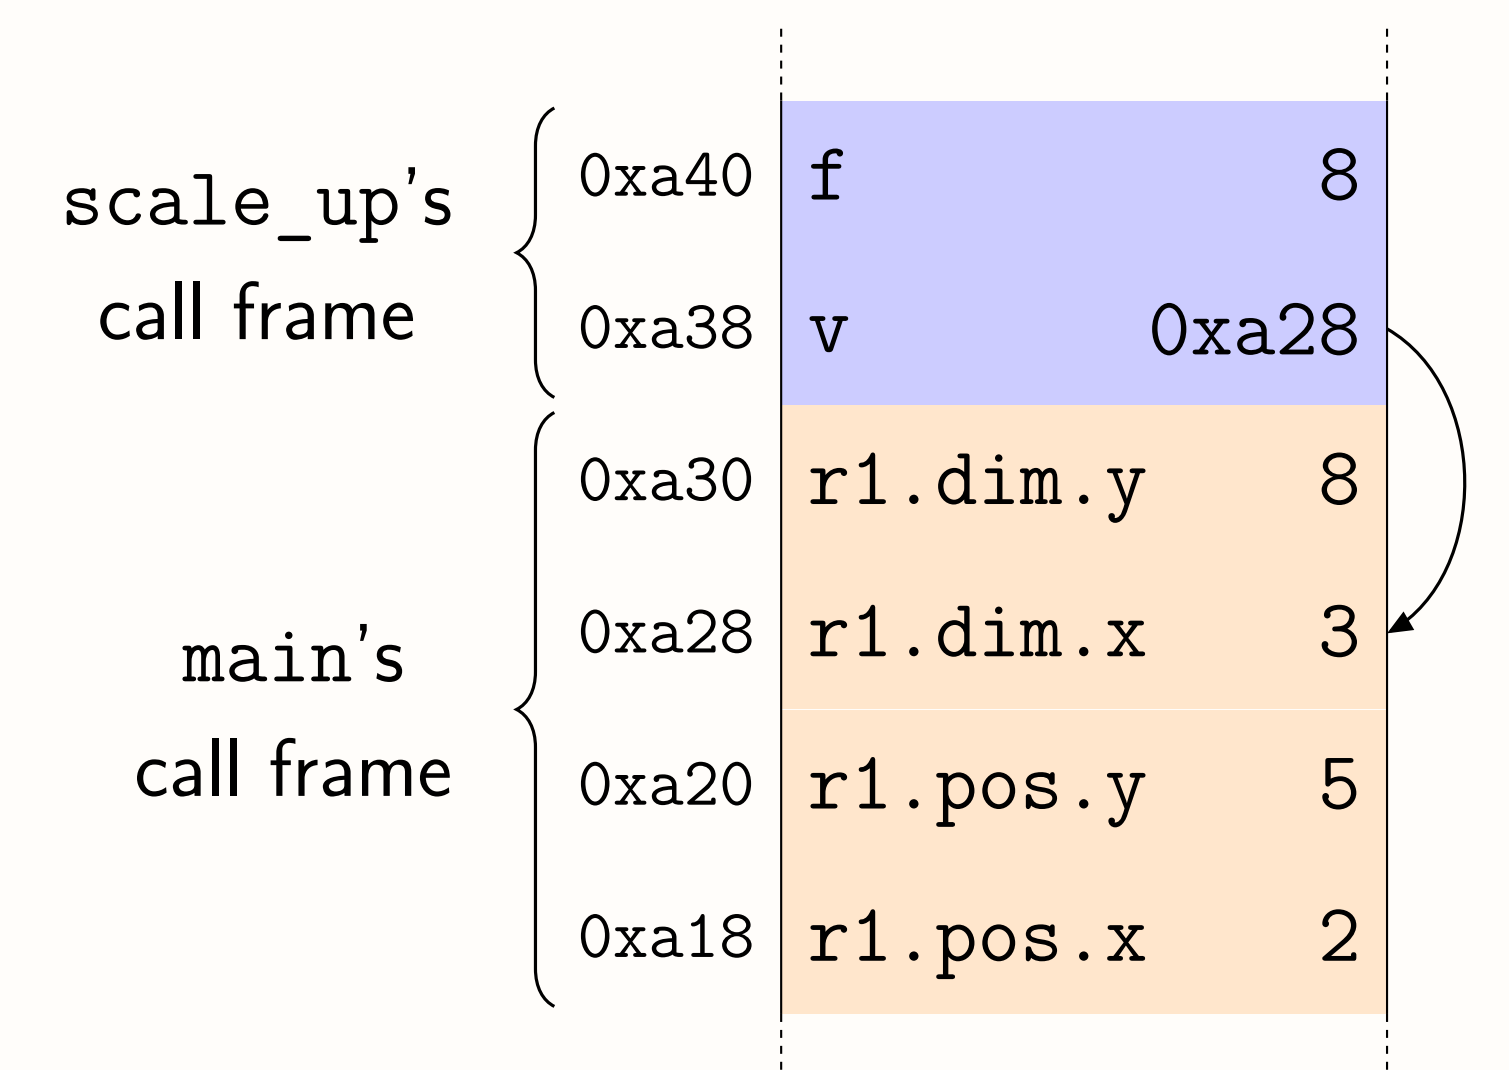
References

- [1] Dimitri Racordon, Denys Shabalin, Daniel Zheng, Dave Abrahams and Brennan Saeta. 2021. Native Implementation of Mutable Value Semantics. *ICOOOLPS@ECOOP*
- [2] <https://github.com/kyouko-taiga/mvs-calculus>.
- [3] <https://blog.rust-lang.org/2017/09/05/Rust-2017-Survey-Results.html>

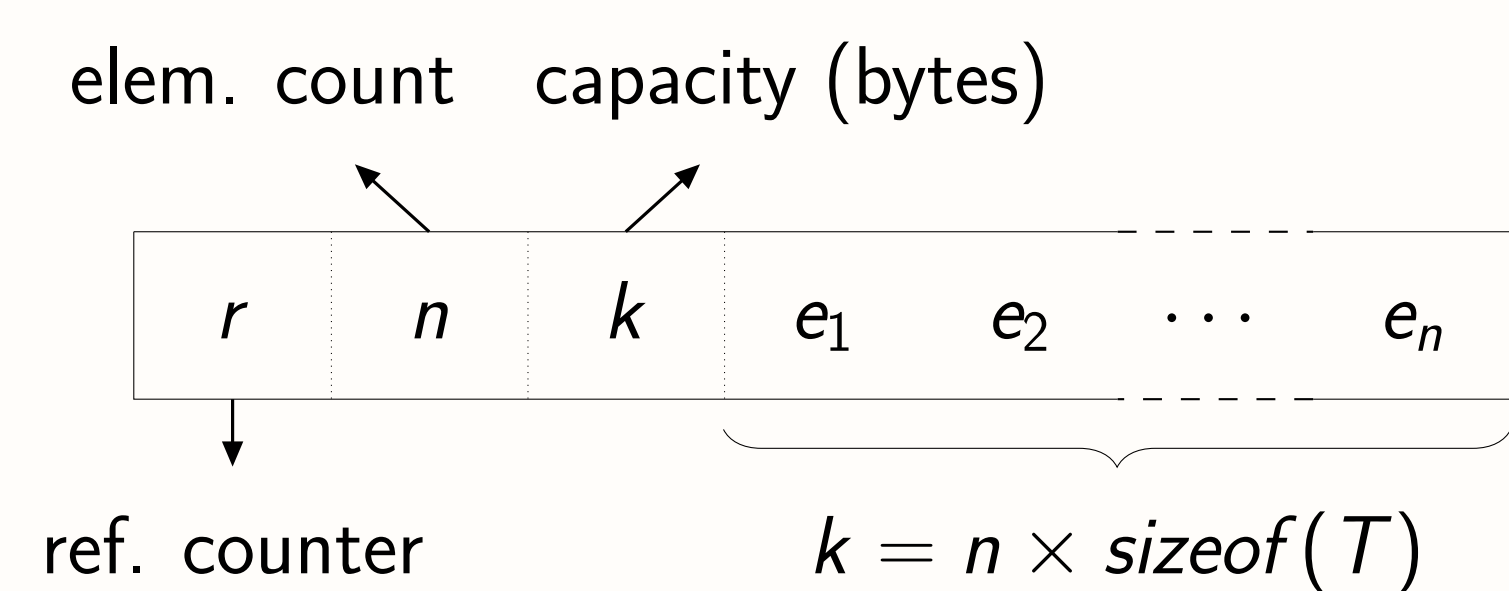
Implementation

Our implementation relies heavily on stack allocation. Built-in types typically have a 1-to-1 correspondence with machine types. Structures (a.k.a. records) are laid out sequentially.

```
var r1 = Rect(Vec2(2,5), Vec2(3,8))
scale_up(&r1.dim, 8)
```



Structures are exploded into scalar arguments and passed directly through registers, **inout** arguments are passed as (possibly interior) pointers.



An array is a pointer to a block of heap-allocated memory that is reference-counted.

Optimizations

Move semantics: Values are often assigned just after having been created (e.g., `var x = [1, 2]`). In this situation, one can *move* the temporary value into the variable and avoid unnecessary heap allocation.

Copy-on-write: Array storage is shared and copied only when mutation occurs, using a reference counter to keep track of sharing.

Dead store elimination: The optimizer leverages local reasoning to discard unnecessary mutations.

Copy propagation: The optimizer identifies immutable data structures and eliminates redundant copies, or substitutes them with aliases to avoid unnecessary allocations.

Evaluation

We measured the execution time of randomly generated programs (250 samples), compiled with our implementation [1, 2], and compared our results with equivalent programs in Swift and C++.

